

Customer Retention Prediction

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, classification_report, roc_auc_score
```

Uploading the dataset

```
[2]: df = pd.read_excel(r"/content/drive/MyDrive/Telco_customer_churn (1).xlsx")
```

```
[3]: # Display First Few Rows
print(df.head())
```

	CustomerID	Count	Country	State	City	Zip Code
0	3668-QPYBK	1	United States	California	Los Angeles	90003
1	9237-HQITU	1	United States	California	Los Angeles	90005
2	9305-CDSKC	1	United States	California	Los Angeles	90006
3	7892-POOKP	1	United States	California	Los Angeles	90010
4	0280-XJGEX	1	United States	California	Los Angeles	90015

	Lat	Long	Latitude	Longitude	Gender	Contract
0	33.964131	-118.272783	33.964131	-118.272783	Male	Month-to-month
1	34.059281	-118.30742	34.059281	-118.307420	Female	Month-to-month
2	34.048013	-118.293953	34.048013	-118.293953	Female	Month-to-month
3	34.062125	-118.315709	34.062125	-118.315709	Female	Month-to-month
4	34.039224	-118.266293	34.039224	-118.266293	Male	Month-to-month

	Paperless Billing	Payment Method	Monthly Charges	Total Charges
0	Yes	Mailed check	53.85	108.15
1	Yes	Electronic check	70.70	151.65
2	Yes	Electronic check	99.65	820.5
3	Yes	Electronic check	104.80	3046.05
4	Yes	Bank transfer (automatic)	103.70	5036.3

	Churn Label	Churn Value	Churn Score	CLTV	Churn Reason
0	Yes	1	86	3239	Competitor made better offer
1	Yes	1	67	2701	Moved
2	Yes	1	86	5372	Moved
3	Yes	1	84	5003	Moved
4	Yes	1	89	5340	Competitor had better devices

[5 rows x 33 columns]

[4]: Data Preprocessing and Handling Missing Values

```
numeric_df = df.select_dtypes(include=np.number)
```

[5]: `df[numeric_df.columns] = df[numeric_df.columns].fillna(numeric_df.median())`

```
# Encoding Categorical Variables
```

```
label_enc = LabelEncoder()
```

```
categorical_cols = df.select_dtypes(include=['object']).columns
```

```
for col in categorical_cols:
```

```
    if df[col].apply(type).nunique() > 1:
```

```
        df[col] = df[col].astype(str)
```

[6]: `df[col] = label_enc.fit_transform(df[col])`

```
# Feature Selection
```

```
features = [ 'Monthly Charges', 'Total Charges', 'CLTV', 'Churn Score', 'Payment_
Method', 'Contract']
```

```
target = 'Churn Value'
```

[7]:

```
X = df[features]
```

[8]:

```
y = df[target]
```

[9]:

```
y = label_enc.fit_transform(y)
```

```
# Train-Test Split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
```

[10]:

```
random_state=42)
```

```
scaler = StandardScaler()
```

```
X_train = scaler.fit_transform(X_train)
```

[11]:

```
!pip install --upgrade scikit-learn xgboost
```

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/dist-packages (1.1.3)

Collecting scikit-learn

Downloading scikit_learn-1.6.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (18 kB)

Requirement already satisfied: xgboost in /usr/local/lib/python3.11/dist-

packages (2.1.4)

Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.26.4)

Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.13.1)

Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.4.2)

Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.5.0)

Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.21.5)

Downloading

scikit_learn-1.6.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
(13.5 MB)

13.5/13.5 MB

51.4 MB/s eta 0:00:00

Installing collected packages: scikit-learn

Attempting uninstall: scikit-learn

Found existing installation: scikit-learn 1.1.3

Uninstalling scikit-learn-1.1.3:

Successfully uninstalled scikit-learn-1.1.3

Successfully installed scikit-learn-1.6.1

```
[12]: import sklearn
import xgboost

print("Scikit-learn version:", sklearn.__version__)
print("XGBoost version:", xgboost.__version__)
```

Scikit-learn version: 1.1.3

XGBoost version: 2.1.4

```
[13]: !pip install scikit-learn==1.1.3
```

Collecting scikit-learn==1.1.3

Using cached scikit_learn-1.1.3-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (10 kB)

Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.11/dist-packages (from scikit-learn==1.1.3) (1.26.4)

Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.11/dist-packages (from scikit-learn==1.1.3) (1.13.1)

Requirement already satisfied: joblib>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn==1.1.3) (1.4.2)

Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn==1.1.3) (3.5.0)

Using cached

scikit_learn-1.1.3-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
(32.0 MB)

Installing collected packages: scikit-learn

Attempting uninstall: scikit-learn

Found existing installation: scikit-learn 1.6.1

Uninstalling scikit-learn-1.6.1:

Successfully uninstalled scikit-learn-1.6.1

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.

mlxtend 0.23.4 requires scikit-learn>=1.3.1, but you have scikit-learn 1.1.3 which is incompatible.

sklearn-compat 0.1.3 requires scikit-learn<1.7,>=1.2, but you have scikit-learn 1.1.3 which is incompatible.

imbalanced-learn 0.13.0 requires scikit-learn<2,>=1.3.2, but you have scikit-learn 1.1.3 which is incompatible.

bigframes 1.34.0 requires scikit-learn>=1.2.2, but you have scikit-learn 1.1.3 which is incompatible.

Successfully installed scikit-learn-1.1.3

Model Training

```
[27]: model = XGBClassifier()  
      model.fit(X_train, y_train)
```

```
[27]: XGBClassifier(base_score=None, booster=None, callbacks=None,  
                  colsample_bylevel=None, colsample_bynode=None,  
                  colsample_bytree=None, device=None, early_stopping_rounds=None,  
                  enable_categorical=False, eval_metric=None, feature_types=None,  
                  gamma=None, grow_policy=None, importance_type=None,  
                  interaction_constraints=None, learning_rate=None, max_bin=None,  
                  max_cat_threshold=None, max_cat_to_onehot=None,  
                  max_delta_step=None, max_depth=None, max_leaves=None,  
                  min_child_weight=None, missing=nan, monotone_constraints=None,  
                  multi_strategy=None, n_estimators=None, n_jobs=None,  
                  num_parallel_tree=None, random_state=None, ...)
```

```
[15]: # Predictions  
      y_pred = model.predict(X_test)
```

Model Evaluation

```
[16]: accuracy = accuracy_score(y_test, y_pred)  
      roc_score = roc_auc_score(y_test, y_pred)  
      print("Accuracy:", accuracy)
```

```
print("ROC-AUC Score:", roc_score)
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9112845990063875


ROC-AUC Score: 0.8852452923686819

	precision	recall	f1-score	support
0	0.93	0.95	0.94	1009
1	0.86	0.82	0.84	400
accuracy			0.91	1409
macro avg	0.89	0.89	0.89	1409
weighted avg	0.91	0.91	0.91	1409

Feature Importance

[17]: # Feature Importance

```
feature_importance = model.feature_importances_
plt.figure(figsize=(8,6))
sns.barplot(x=features, y=feature_importance)
plt.xlabel("Feature")
plt.ylabel("Importance")
plt.title("Feature Importance in Churn Prediction")
plt.show()
```



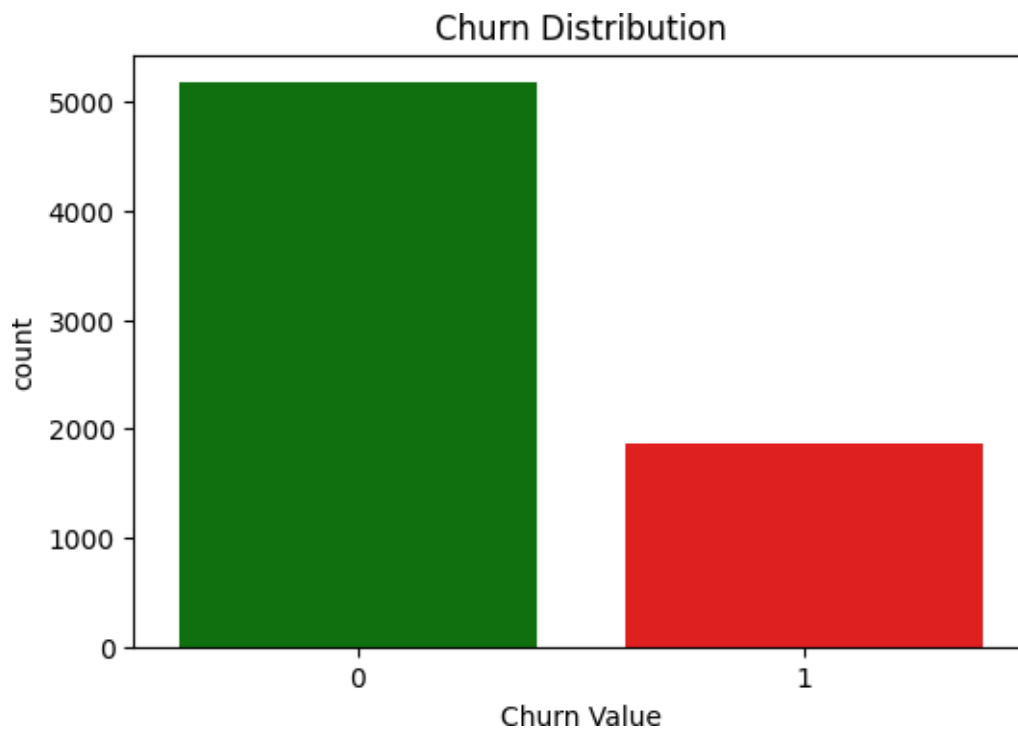
Customer Retention & Churn Overview

```
[19]: plt.figure(figsize=(6,4))  
sns.countplot(x='Churn Value', data=df, palette=['green', 'red'])  
plt.title('Churn Distribution')  
plt.show()
```

<ipython-input-19-559b51ffd2cb>:6: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Churn Value', data=df, palette=['green', 'red'])
```



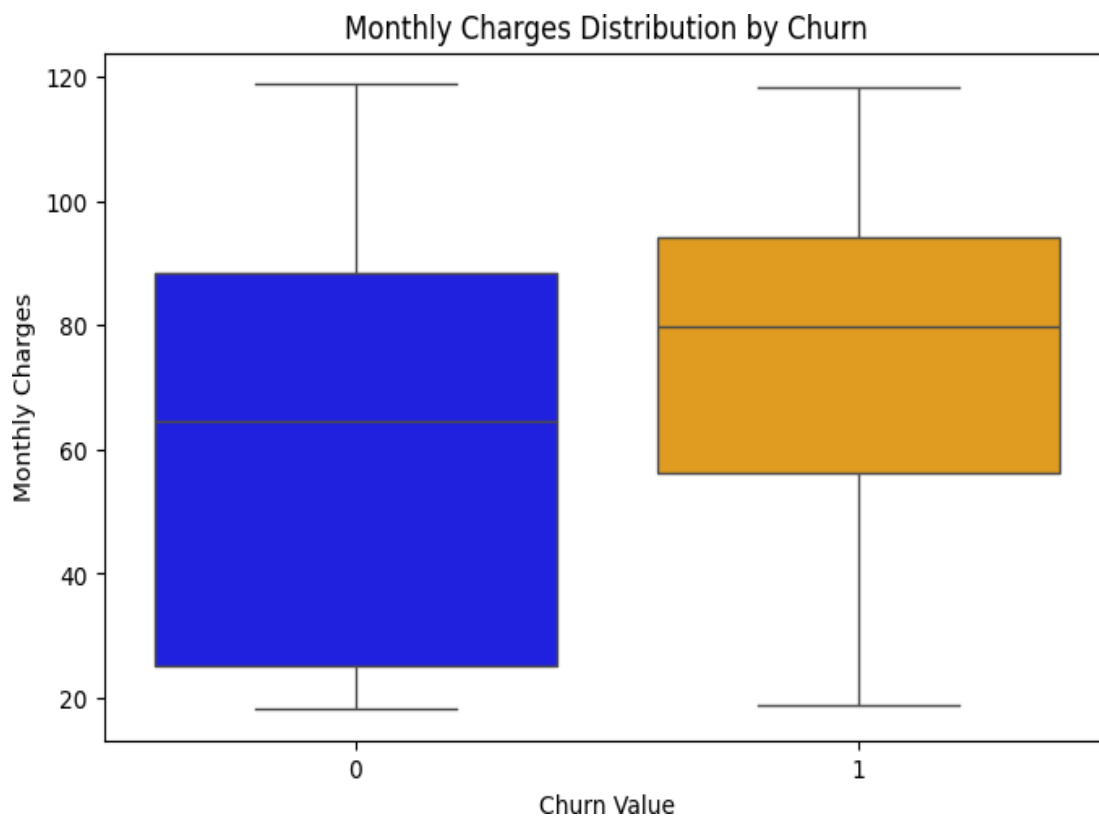
Monthly Charges vs. Churn

```
[24]: plt.figure(figsize=(8,5))
sns.boxplot(x='Churn Value', y='Monthly Charges', data=df, palette=['blue',
    'orange'])
plt.title('Monthly Charges Distribution by Churn')
plt.show()
```

<ipython-input-24-b61ab3eafaf2>:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Churn Value', y='Monthly Charges', data=df, palette=['blue',
'orange'])
```



Model Performance Evaluation

```
[26]: y_pred = model.predict(X_test)
      cm = confusion_matrix(y_test, y_pred)

      plt.figure(figsize=(6,4))
      sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
      plt.xlabel('Predicted')
      plt.ylabel('Actual')
      plt.title('Confusion Matrix')
      plt.show()
```

